

THEORY



Training Program 01

Advanced NLP & Large Language Models (LLMs)



CONFIANZA23
Inteligencia y Seguridad

M01: Text Preprocessing

M02: Word and Sentence Embeddings

M03: Transformer Architectures

M04: Entity-Centric NLP

M05: Applied NLP Tasks

M06: Multimodal NLP

M07: LLM Fundamentals

M08: Prompt Engineering

M09: Fine-Tuning & PEFT

M10: RAG & Real-World Projects



MODULE M01

TEXT PREPROCESSING

CHAPTER M05

Applied NLP Tasks



CONFIANZA23
Inteligencia y Seguridad

5.1 Introduction to Applied NLP Tasks

5.2 Core Task: Text Classification

5.3 Understanding User Intent: Sentiment Analysis

5.4 Reducing Information Overload: Text Summarization

5.5 Information Retrieval: Question Answering Systems


5.6 Task Comparison and the Future of Applied NLP



Chapter 5.1 Introduction to Applied NLP Tasks

From Theory to Applied NLP


- NLP moves from mathematical modeling to real-world problem solving.
- Applied NLP focuses on business, industrial, and societal value.
- Models are evaluated by usefulness, not only benchmark scores.
- This module connects Transformers to production systems.



Chapter 5.1 Introduction to Applied NLP Tasks

Applied NLP Task Families


- Applied NLP tasks fall into two paradigms
- Tasks are defined by what the model produces
- Choice impacts architecture, evaluation, and cost



Chapter 5.1 Introduction to Applied NLP Tasks

Discriminative Tasks

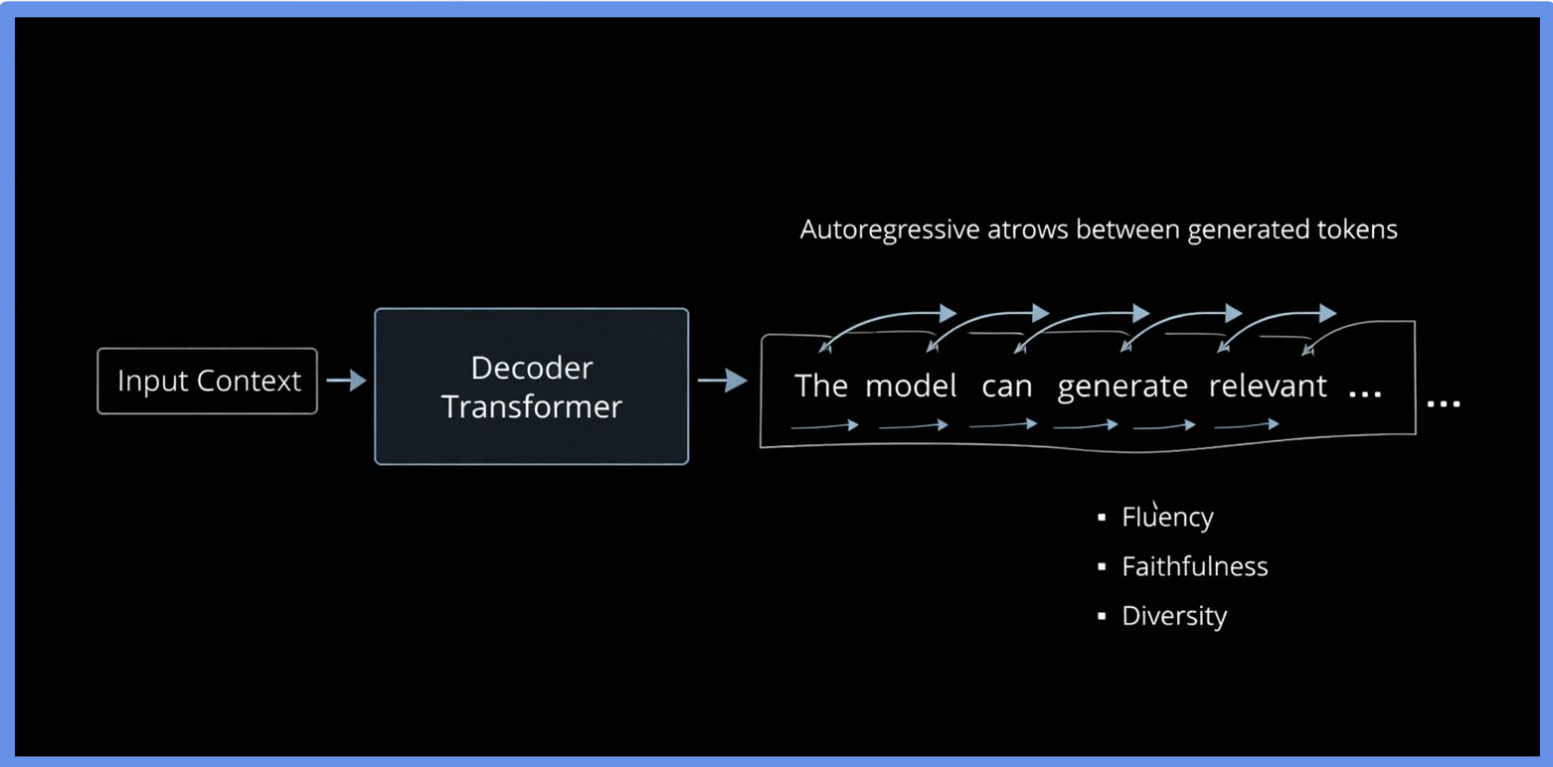
- Output: predefined labels
- Goal: decision, filtering, categorization
- Typical architectures: Encoder-only (BERT-style)
- Metrics: Accuracy, Precision, Recall, F1-score
- Examples: classification, sentiment, NER



Chapter 5.1 Introduction to Applied NLP Tasks

Generative Tasks


- Output: newly generated text
- Goal: synthesis, explanation, abstraction
- Typical architectures: Decoder or Encoder–Decoder
- Evaluation is probabilistic and semantic
- Examples: summarization, QA, text generation



Chapter 5.1 Introduction to Applied NLP Tasks

Benchmarks and Model Readiness


- Benchmarks standardize model comparison
- GLUE and SuperGLUE measure reasoning ability
- High benchmark scores $\neq$ production readiness
- Fine-tuning adapts models to domain-specific tasks



Chapter 5.1 Introduction to Applied NLP Tasks

Production Challenges in Applied NLP

- Benchmarks standardize model comparison.
- GLUE and SuperGLUE measure reasoning ability.
- High benchmark scores $\neq$ production readiness.
- Fine-tuning adapts models to domain-specific tasks.



Chapter 5.1 Introduction to Applied NLP Tasks

Modular Pipelines vs End-to-End Models

- End-to-End: simpler, faster to prototype
- Modular pipelines: more explainable and controllable
- Enterprises often combine both approaches
- Design choice depends on scale and regulation

Monolithic vs. Modular LLM Systems


Characteristic	Monolithic LLM	Modular NLP Pipeline
Structure	Single large rectangular box	Multiple rectangular boxes connected left-to-right
Internal Detail	Minimal, just the label	NER, Classifier, Retriever, Generator
Flow	No arrows needed	NER → Classifier → Retriever → Generator
Comparison	Contrasts with multi-module pipeline	Contrasts with single-box system
Style	Academic / engineering blueprint	Academic / engineering blueprint
Palette	Monochrome (white and gray) on black	Monochrome (white and gray) on black
Accent Color	Optional (light blue or teal)	Optional (light blue or teal)
Typography	Sans-serif, neutral	Sans-serif, neutral
Design	Flat design	Flat design
Alignment	Horizontally aligned	Horizontally aligned
Intent	Visual comparison	Visual comparison



Chapter 5.1 Introduction to Applied NLP Tasks

Why This Module Matters

- Applied NLP turns language into action.
- Tasks define how AI creates value.
- Understanding task structure prevents misuse.
- Next chapters deep-dive into each task.



5.1 Introduction to Applied NLP Tasks

5.2 Core Task: Text Classification

5.3 Understanding User Intent: Sentiment Analysis

5.4 Reducing Information Overload: Text Summarization

5.5 Information Retrieval: Question Answering Systems


5.6 Task Comparison and the Future of Applied NLP



Chapter 5.2 Text Classification

What is Text Classification?

- Task: assign predefined labels to text.
- One of the oldest and most deployed NLP tasks.
- Foundation for moderation, routing, filtering, analytics.
- Used in almost every production NLP system.



Chapter 5.2 Text Classification

Classification Problem

Types

- Binary: two exclusive classes (Yes / No)
- Multi-class: one label from many
- Multi-label: multiple labels per document
- Correct framing defines model architecture


Classification Paradigms				
Characteristic	 Output Boxes	 Active Outputs	 Arrow	 Title
Binary Classification	Two	One	Single	Binary Classification
Multi-class Classification	Multiple	One	To set	Multi-class Classification
Multi-label Classification	Multiple	Several	To set	Multi-label Classification



Chapter 5.2 Text Classification

Traditional Classification Methods

- Based on word frequency statistics.
- Require manual feature engineering.
- Still useful for small or structured datasets.
- Serve as strong baselines.




Chapter 5.2 Text Classification

Deep Learning Classification

Models learn features automatically.


- Capture semantic meaning, not just keywords.
- Transformers are the modern standard.
- Context-aware classification.



Chapter 5.2 Text Classification

Zero-Shot Classification


- No task-specific training required.
- Uses natural language label descriptions.
- Extremely fast to deploy.
- Ideal for dynamic taxonomies.



Chapter 5.2 Text Classification

Few-Shot Classification


- No task-specific training required
- Uses natural language label descriptions
- Extremely fast to deploy
- Ideal for dynamic taxonomies



Chapter 5.2 Text Classification

Evaluation Metrics


- Accuracy is often misleading
- Precision controls false positives
- Recall controls false negatives
- F1-score balances both
- Metric choice depends on risk



Chapter 5.2 Text Classification

Key Takeaways

- Classification transforms language into decisions
- Architecture choice matters more than model size
- Zero-shot enables rapid deployment
- Metrics define real-world success



5.1 Introduction to Applied NLP Tasks

5.2 Core Task: Text Classification

5.3 Understanding User Intent: Sentiment Analysis

5.4 Reducing Information Overload: Text Summarization

5.5 Information Retrieval: Question Answering Systems


5.6 Task Comparison and the Future of Applied NLP



Chapter 5.3 Understanding User Intent: Sentiment Analysis

What is Sentiment Analysis?


- Task: detect subjective intent in text
- Focuses on opinions, emotions, attitudes
- Specialized form of text classification
- Core signal for user and market intelligence



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Document-Level Sentiment


- One sentiment for the entire document
- Simple and computationally efficient
- Works well for short reviews or posts
- Loses internal nuance



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Document-Level Sentiment


- One sentiment for the entire document
- Simple and computationally efficient
- Works well for short reviews or posts
- Loses internal nuance



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Aspect-Based Sentiment Analysis (ABSA)


- Sentiment linked to specific entities or features
- Enables fine-grained insight
- Essential for product and brand analysis
- Industry standard for actionable feedback



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Lexicon-Based Approaches


- Use predefined polarity dictionaries
- Fully explainable
- Fast and lightweight
- Fail under context, irony, or negation



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Embedding-Based Sentiment Models


- Meaning represented in latent vector space
- Generalizes beyond exact keywords
- Robust to synonyms and paraphrases
- Foundation for transformer-based sentiment



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Negation and Sarcasm Challenges

- Negation flips polarity meaning
- Sarcasm contradicts surface wording
- Requires long-range context
- Transformers outperform classical models



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Sentiment in Market Intelligence


- Meaning represented in latent vector space
- Generalizes beyond exact keywords
- Robust to synonyms and paraphrases
- Foundation for transformer-based sentiment



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Operational Considerations


- Domain-specific language matters
- Models degrade with cultural drift
- Requires periodic re-evaluation
- Sentiment $\neq$ truth



Chapter 5.3 Understanding User Intent: Sentiment Analysis

Key Takeaways

- Sentiment reveals user intent.
- Aspect-level analysis enables action.
- Context is the main technical challenge.
- Critical signal for applied NLP systems.



5.1 Introduction to Applied NLP Tasks

5.2 Core Task: Text Classification

5.3 Understanding User Intent: Sentiment Analysis

5.4 Reducing Information Overload: Text Summarization

5.5 Information Retrieval: Question Answering Systems


5.6 Task Comparison and the Future of Applied NLP



Chapter 5.4 Reducing Information Overload: Text Summarization

What is Text Summarization?


- Task: compress long text into key information.
- Preserves meaning while reducing length.
- Essential for scale and decision speed.
- Widely used in legal, news, intelligence.



Chapter 5.4 Reducing Information Overload: Text Summarization

Extractive Summarization


- Task: compress long text into key information
- Preserves meaning while reducing length
- Essential for scale and decision speed
- Widely used in legal, news, intelligence



Chapter 5.4 Reducing Information Overload: Text Summarization

Abstractive Summarization


- Generates new sentences
- Higher information density
- More natural language
- Risk of hallucination



Chapter 5.4 Reducing Information Overload: Text Summarization

Traditional Extractive Techniques


- Graph-based ranking (TextRank)
- Sentence similarity scoring
- No deep semantics
- Still useful for formal text



Chapter 5.4 Reducing Information Overload: Text Summarization

Transformer-Based Summarization

- Uses full-document context
- Learns semantic importance
- State-of-the-art performance
- Dominant in production systems




M05

Chapter 5.4 Reducing Information Overload: Text Summarization

BART and T5 Models

- UEncoder–Decoder architectures
- Pre-trained for text reconstruction
- Strong performance on long documents
- Industry standard for summarization




CONFIANZA23
Inteligencia y Seguridad

Chapter 5.4 Reducing Information Overload: Text Summarization

Evaluation Metrics (ROUGE)

- Measures overlap with human summaries
- ROUGE-1, ROUGE-2, ROUGE-L
- Focuses on recall
- Does not ensure factuality




M05

Chapter 5.4 Reducing Information Overload: Text Summarization

Key Takeaways

- Summarization manages information overload
- Extractive = safer, abstractive = denser
- Transformers dominate modern systems



CONFIANZA23
Inteligencia y Seguridad

5.1 Introduction to Applied NLP Tasks

5.2 Core Task: Text Classification

5.3 Understanding User Intent: Sentiment Analysis

5.4 Reducing Information Overload: Text Summarization

5.5 Information Retrieval: Question Answering

Systems


5.6 Task Comparison and the Future of Applied NLP



Chapter 5.5 Information Retrieval: Question Answering Systems

What is Question Answering (QA)?


- Goes beyond keyword search
- Requires language understanding
- Core of modern AI assistants



Chapter 5.5 Information Retrieval: Question Answering Systems

Extractive Question Answering

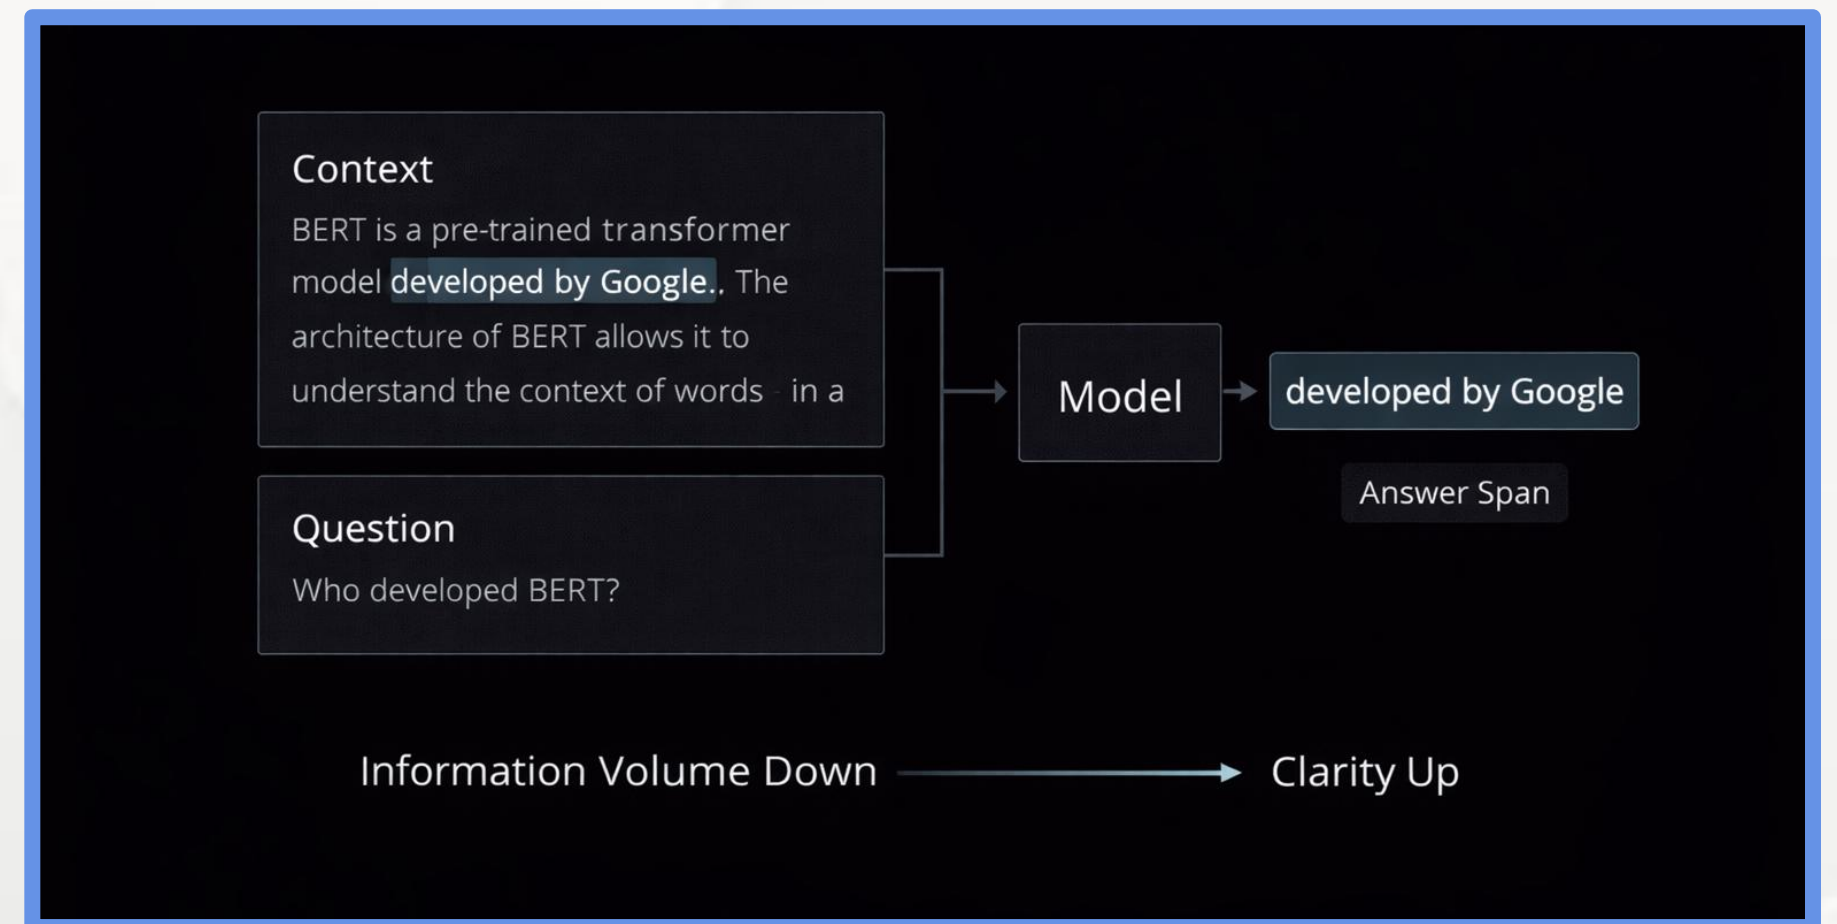
- Answer is a text span
- Start and end token prediction
- High factual precision
- Limited to provided context



Chapter 5.5 Information Retrieval: Question Answering Systems

The SQuAD Paradigm


- Benchmark dataset for extractive QA
- Context + Question → Span
- Enabled large-scale QA training
- Foundation for enterprise FAQ systems



Chapter 5.5 Information Retrieval: Question Answering Systems

Generative Question Answering


- Benchmark dataset for extractive QA
- Context + Question → Span
- Enabled large-scale QA training
- Foundation for enterprise FAQ systems



Chapter 5.5 Information Retrieval: Question Answering Systems

Knowledge-Based Question Answering


- Uses structured databases
- Converts language to queries
- Deterministic answers
- High reliability systems



Chapter 5.5 Information Retrieval: Question Answering Systems

Retrieval-Augmented Generation (RAG)


- Industry-standard QA architecture
- Combines retrieval + generation
- Uses private documents
- Reduces hallucinations



Chapter 5.5 Information Retrieval: Question Answering Systems

Vector Databases in QA

- Exact Match (EM)
- F1-Score (span overlap)
- Faithfulness (RAG-specific)
- Citation correctness




M05

Chapter 5.5 Information Retrieval: Question Answering Systems

Evaluation Metrics for QA

- Customer support automation
- Technical documentation search
- Legal and medical assistants
- Enterprise knowledge access




CONFIANZA23
Inteligencia y Seguridad

Chapter 5.5 Information Retrieval: Question Answering Systems

Key Takeaways

- Customer support automation
- Technical documentation search
- Legal and medical assistants
- Enterprise knowledge access



5.1 Introduction to Applied NLP Tasks

5.2 Core Task: Text Classification

5.3 Understanding User Intent: Sentiment Analysis

5.4 Reducing Information Overload: Text Summarization

5.5 Information Retrieval: Question Answering Systems


5.6 Task Comparison and the Future of Applied NLP



Chapter 5.6 Task Comparison and the Future of Applied NLP

What is Question Answering?

- Customer support automation
- Technical documentation search
- Legal and medical assistants
- Enterprise knowledge access




Chapter 5.6 Task Comparison and the Future of Applied NLP

Extractive vs Generative QA

QA systems operate under two main paradigms: extracting exact answers from text or generating new answers using language models.

- Extractive: answer is a text span
- High factual precision
- Used in legal and technical domains
- Generative: produces new text
- More fluent, higher hallucination risk




Chapter 5.6 Task Comparison and the Future of Applied NLP

Retrieval-Augmented Generation (RAG)

Retrieval-Augmented Generation (RAG) combines document retrieval with generative models to produce answers grounded in real data.

- Industry standard for QA
- Uses vector search
- Injects retrieved text into prompts
- Reduces hallucinations
- Enables private data QA




Chapter 5.6 Task Comparison and the Future of Applied NLP

Evaluation and Production Use

Reliable QA systems must be evaluated not only for correctness, but also for trust and factual grounding.

- Exact Match (EM)
- F1-score (span overlap)
- Faithfulness (RAG answers)
- Enterprise QA systems
- Trust and explainability



Resources & Academic References

Core Benchmarks & Tasks

Wang et al. (2018) — GLUE: A Multi-Task Benchmark for Natural Language Understanding
<https://arxiv.org/abs/1804.07461>

Wang et al. (2019) — SuperGLUE: A Stickier Benchmark for General-Purpose Language Understanding Systems
<https://arxiv.org/abs/1905.00537>

Rajpurkar et al. (2016) — SQuAD: 100,000+ Questions for Machine Comprehension of Text
<https://arxiv.org/abs/1606.05250>

Text Classification & Zero-Shot Learning

Cortes & Vapnik (1995) — Support-Vector Networks
<https://link.springer.com/article/10.1007/BF00994018>

Yin et al. (2019) — Benchmarking Zero-shot Text Classification
<https://arxiv.org/abs/1909.00161>



Resources & Academic References

Sentiment Analysis

Pang & Lee (2008) — Opinion Mining and Sentiment Analysis
<https://www.cs.cornell.edu/home/llee/omsa/omsa.pdf>

Pontiki et al. (2014) — SemEval-2014 Task 4: Aspect-Based Sentiment Analysis
<https://aclanthology.org/S14-2004.pdf>

Text Summarization

Barzilay & Elhadad (1997) — Using Lexical Chains for Text Summarization
<https://aclanthology.org/P97-1002.pdf>

Lewis et al. (2019) — BART: Denoising Sequence-to-Sequence Pre-training
<https://arxiv.org/abs/1910.13461>

Raffel et al. (2019) — Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer (T5)
<https://arxiv.org/abs/1910.10683>



Resources & Academic References

■ Question Answering & RAG

Lewis et al. (2020) — Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks

<https://arxiv.org/abs/2005.11401>

■ Ethics & Future of Applied NLP

Bender et al. (2021) — On the Dangers of Stochastic Parrots

<https://dl.acm.org/doi/pdf/10.1145/3442188.3445922>

Bommasani et al. (2021) — On the Opportunities and Risks of Foundation Models

<https://arxiv.org/abs/2108.07258>



CONFIANZA23
Inteligencia y Seguridad