

THEORY



Training Program 01

Advanced NLP & Large Language Models (LLMs)



CONFIANZA23
Inteligencia y Seguridad

Training Program 01

M01: Text Preprocessing

M02: Word and Sentence Embeddings

M03: Transformer Architectures

M04: Entity-Centric NLP

M05: Applied NLP Tasks

M06: Multimodal NLP

M07: LLM Fundamentals

M08: Prompt Engineering

M09: Fine-Tuning & PEFT

M10: RAG & Real-World Projects



MODULE M01

TEXT PREPROCESSING

CHAPTER M09

Fine-Tuning & PEFT



CONFIANZA23
Inteligencia y Seguridad

9.1 The Need for Efficient Fine-Tuning

9.2 LoRA: Low-Rank Adaptation Mechanics

9.3 QLoRA: Quantization for Parameter Efficiency

9.4 RLHF: Aligning Models with Human Feedback

9.5 Practical PEFT: A Hands-on Fine-Tuning Project

9.6 Summary and Next Steps




Chapter 9.1 The Need for Efficient Fine-Tuning

Why Full Fine-Tuning No Longer Scales

Modern LLMs are too large to be fully fine-tuned in practical environments. Economic and hardware constraints force a shift toward efficiency-first adaptation strategies.

- Full fine-tuning updates every parameter
- 70B models exceed realistic VRAM limits
- Costs scale linearly with model size
- Not viable for most organizations




Chapter 9.1 The Need for Efficient Fine-Tuning

The VRAM Bottleneck

Training requires far more memory than inference due to gradients, optimizer states, and activations.

- Full fine-tuning updates every parameter
- 70B models exceed realistic VRAM limits
- Costs scale linearly with model size
- Not viable for most organizations




Chapter 9.1 The Need for Efficient Fine-Tuning

Catastrophic Forgetting

Aggressive fine-tuning on narrow datasets can erase general reasoning abilities.

- Model overwrites prior knowledge
- Loss of instruction-following skills
- Reduced generalization
- Especially risky with small datasets




Chapter 9.1 The Need for Efficient Fine-Tuning

Why PEFT Became the Standard

PEFT freezes the base model and trains only a small number of new parameters.

- 90–95% VRAM reduction
- Preserves general intelligence
- Enables fast experimentation
- Ideal for domain adaptation



Training Program 01

9.1 The Need for Efficient Fine-Tuning

9.2 LoRA: Low-Rank Adaptation Mechanics

9.3 QLoRA: Quantization for Parameter Efficiency

9.4 RLHF: Aligning Models with Human Feedback

9.5 Practical PEFT: A Hands-on Fine-Tuning Project

9.6 Summary and Next Steps




Chapter 9.2 LoRA: Low-Rank Adaptation Mechanics

The Core Idea of LoRA

LoRA assumes task-specific updates live in a low-dimensional subspace.

- Weight updates are low-rank
- No need to modify full matrices
- Efficient and expressive
- Minimal disruption to base model




Chapter 9.2 LoRA: Low-Rank Adaptation Mechanics

The Core Idea of LoRA

Instead of updating W , LoRA adds a parallel low-rank path.

- W remains frozen
- $\Delta W = B \cdot A$
- Bottleneck enforces efficiency
- Adapter path merged at inference



Chapter 9.2 LoRA: Low-Rank Adaptation Mechanics

Rank and Alpha Trade-offs

LoRA performance depends on careful hyperparameter selection.

- Rank controls capacity
- Alpha controls update strength
- Higher $\neq$ always better
- Typical r : 8–32

Comparison Chart: Full Fine-Tuning vs LoRA

	Full Fine-Tuning	LoRA
Parameters	All parameters	Low-rank updates
Storage	High storage cost	Low storage cost
Speed	 Slow	 Faster

LoRA adapter path

Adapter path merged
at inference



Training Program 01

9.1 The Need for Efficient Fine-Tuning

9.2 LoRA: Low-Rank Adaptation Mechanics

9.3 QLoRA: Quantization for Parameter Efficiency

9.4 RLHF: Aligning Models with Human Feedback

9.5 Practical PEFT: A Hands-on Fine-Tuning Project

9.6 Summary and Next Steps




Chapter 9.3 QLoRA: Quantization for Parameter Efficiency

Why Quantization Is Necessary

Even frozen models are too large to load in FP16.

- 70B FP16 $\approx$ 140 GB VRAM
- Consumer GPUs cannot load them
- Quantization solves the loading problem




Chapter 9.3 QLoRA: Quantization for Parameter Efficiency

QLoRA: Training Through 4-bit Weights

QLoRA combines 4-bit quantization with LoRA adapters.

- Base model in NF4 (4-bit)
- LoRA trained in higher precision
- On-the-fly dequantization
- No accuracy collapse




Chapter 9.3 QLoRA: Quantization for Parameter Efficiency

Double Quantization

QLoRA reduces memory even further by quantizing quantization constants.

- Saves several GB VRAM
- Critical for large models
- Enables single-GPU fine-tuning



Training Program 01

9.1 The Need for Efficient Fine-Tuning

9.2 LoRA: Low-Rank Adaptation Mechanics

9.3 QLoRA: Quantization for Parameter Efficiency

9.4 RLHF: Aligning Models with Human Feedback

9.5 Practical PEFT: A Hands-on Fine-Tuning Project

9.6 Summary and Next Steps




Chapter 9.4 RLHF: Aligning Models with Human Feedback

The Alignment Problem

Instruction-following alone does not guarantee safe or helpful behavior.

- Models inherit internet biases
- May be verbose, unsafe, or misaligned
- Alignment is mandatory for production



Chapter 9.4 RLHF: Aligning Models with Human Feedback

The RLHF Pipeline

RLHF converts human preference into a learnable signal.

- Supervised Fine-Tuning (SFT)
- Reward Model training
- Reinforcement Learning (PPO)
- Human values → reward




Chapter 9.4 RLHF: Aligning Models with Human Feedback

PPO and Modern Alternatives

PPO optimizes rewards while preserving model stability.

- Prevents over-optimization
- Maintains language quality
- DPO simplifies the pipeline
- Fewer moving parts



Training Program 01

9.1 The Need for Efficient Fine-Tuning

9.2 LoRA: Low-Rank Adaptation Mechanics

9.3 QLoRA: Quantization for Parameter Efficiency

9.4 RLHF: Aligning Models with Human Feedback

9.5 Practical PEFT: A Hands-on Fine-Tuning Project

9.6 Summary and Next Steps




Chapter 9.5 Practical PEFT: A Hands-on Fine-Tuning Project

Project Objective

The goal is to build a domain-expert assistant via PEFT.

- Legal audit assistant
- 2k–5k instruction-response pairs
- Strict data schema
- High-quality supervision




Chapter 9.5 Practical PEFT: A Hands-on Fine-Tuning Project

Training Setup

QLoRA + LoRA enables efficient specialization.

- Legal audit assistant
- 2k–5k instruction-response pairs
- Strict data schema
- High-quality supervision




Chapter 9.5 Practical PEFT: A Hands-on Fine-Tuning Project

Evaluation and Deployment

Training is monitored and adapters are merged for deployment.

- Watch validation loss
- Test unseen samples
- Merge LoRA into base
- Produce single deployable model



Training Program 01

9.1 The Need for Efficient Fine-Tuning

9.2 LoRA: Low-Rank Adaptation Mechanics

9.3 QLoRA: Quantization for Parameter Efficiency

9.4 RLHF: Aligning Models with Human Feedback

9.5 Practical PEFT: A Hands-on Fine-Tuning Project

9.6 Summary and Next Steps



Chapter 9.6 Summary and Next Steps

Choosing the Right Adaptation Strategy

Fine-tuning must be treated as a strategic architectural decision, not as a default solution. Each adaptation technique introduces increasing levels of complexity, cost, and maintenance overhead. A mature AI system follows a layered escalation path, applying the minimum level of intervention required to achieve the desired performance and reliability.

- Prompting first: Use prompt engineering to shape behavior without modifying weights
- RAG second: Inject external knowledge before altering the model itself
- PEFT when necessary: Apply LoRA / QLoRA when style, structure, or domain expertise must be permanent
- Full fine-tuning rarely justified: Reserved for foundation-model creators or extreme specialization scenarios



Chapter 9.6 Summary and Next Steps

The Future of PEFT

Model adaptation is evolving toward dynamic, modular intelligence, where specialization is no longer a static training event but a real-time, routable capability. Instead of maintaining multiple full models, future systems will rely on a single foundation model enhanced by a library of lightweight, interchangeable adapters.

- Adapter routing: Automatic selection of the most relevant LoRA at inference time
- Mixture-of-Experts adapters: Multiple specialized adapters activated in parallel
- Instant specialization: No retraining required per request
- Shared backbone intelligence: One model, many expert behaviors



Key References

- Houlsby et al. (2019) – Parameter-Efficient Transfer Learning for NLP
- Lester et al. (2021) – The Power of Scale for Parameter-Efficient Prompt Tuning
- Hu et al. (2021) – LoRA: Low-Rank Adaptation of Large Language Models
- Aghajanyan et al. (2020) – Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning
- Dettmers et al. (2023) – QLoRA: Efficient Finetuning of Quantized LLMs
- Frantar et al. (2022) – GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers
- Ouyang et al. (2022) – Training Language Models to Follow Instructions with Human Feedback
- Rafailov et al. (2023) – Direct Preference Optimization
- Taori et al. (2023) – Stanford Alpaca: An Instruction-Following LLaMA Model
- Mangrulkar et al. (2022) – PEFT: State-of-the-Art Parameter-Efficient Fine-Tuning Methods
- Sun et al. (2023) – A Survey of Parameter-Efficient Fine-Tuning in Large Language Models





CONFianza23
Inteligencia y Seguridad